
ULTRA-SORT

EXTREMELY PARALLEL HARDWARE OPTIMIZED SORTING

15-618 FINAL PROJECT REPORT

Dee Dong

ddong1@andrew.cmu.edu

Saatvik Shah

saatviks@andrew.cmu.edu

1 INTRODUCTION

Just as a recap, as part of our 15618 project, we propose to build a production level DSL-style sorting system efficiently utilizing the hardware its running on. Here is what it should be capable of:

1. Leveraging SIMD parallelism for the latest hardware - optimized for the latest AVX512 registers.
2. Cache friendly
3. Efficient in space and time.
4. Generalized: Support all numeric data types.
5. Generalized: Supporting sorting n-ary tuples by different sort-keys.
6. Bonus goal: Multi-core.
7. Bonus goal: NUMA Awareness.
8. Bonus goal: Peloton Integration.

We are happy to report that apart from the additional bonus goals of *NUMA awareness and Peloton Integration* we have achieved all our primary goals and the bonus goal of multicore SIMD sorting. In terms of results we are able to [Dee our best results for serial/parallel here](#). The sections have been arranged as follows: Section II provides a literature survey, Section III provides a background and detailed introduction to SIMD and all the components/operations needed within the main algorithm, Section IV covers our main algorithm approach and the variants we've setup. In Section V, we'll finally motivate and introduce the design of our DSL for SIMD-Sorting. We'll present our benchmarking results and analysis in Section VI. In Section VII we'll conclude our report.

2 LITERATURE REVIEW

Sorting is one of the most important phases in many engineering applications. As part of our project, we have decided to experiment with it keeping in mind the DBMS as an application. There has been a lot of research to exploit the parallelism one can get out of a simple operation as sorting: SIMD intrinsics, multicore and even NUMA-awareness. Prior literatures have some exciting performance results regarding SIMD sorting. C. Kim & Dubey (2009) have suggested that, once hardware provides support for vector instructions that are sufficiently wide (SIMD with 256-bit AVX and wider), sort-merge joins would easily outperform radix-hash joins. C. Balkesen (2013) experimented with AVX2 (256 bit) and found out that radix-hash join is still superior than sort-merge join in many cases, but with the access to wider SIMD this might no longer hold. They also implemented NUMA-aware algorithms in which sort-merge join performed better. H. Inoue (2015) proposed NUMA-aware and cache-friendly sorting algorithm but with AVX2, which also indicates the need for updating the state of art using AVX512.

3 BACKGROUND

3.1 NOTATION

1. **Scalar vs. SIMD sort:** As expected, a scalar sort is one which does not use any SIMD/vectorized instructions while a SIMD sort does.
2. **Horizontal vs. Vertical SIMD operations:** Horizontal operations apply on a SIMD register on its own, while Vertical operations apply to a pair or group of registers.
3. **Run-size or Chunk size:** A Run or chunk represents a group of numbers which are being operated on. A `unit run size` refers to a minimal such size.
4. **Comparison vs. Movement operations:** A SIMD operation has been referred to as a comparison operation if the actual values within the register are being compared. It is a movement operation if there is horizontal or vertical movement of values independent of the magnitude of values.

3.2 USEFUL RESOURCES

In this subsection we have laid out a number of useful resources which we've ended up using throughout our project. This simply serves as a future reference to the authors of this document as well as interested readers.


1. [Introduction to Intel Xeon Phi and AVX512](#)
2. [Excellent Introduction to AVX](#)
3. [AVX Intrinsics Guide](#)
4. [Automatic Sorting Network Design](#)

3.3 BASIC SIMD INSTRUCTIONS

We have used a number of AVX intrinsics to achieve relatively complex operations in a vectorized manner. In this subsection we provide insight into some of the building blocks to these more complex operations.

1. **Min/Max:** Directly takes the minimum or maximum and stores into the output register. Below example shows min/max instructions for 4 double precision numbers

Figure 1: Min-Max



2. **Unpack Lo/Hi:** Unpack-Lo first forms **two groups** of equal numbers, and extracts the lower order numbers from each group sequentially into the output register. Unpack-Hi takes the higher order numbers instead.
3. **Shuffle:** The control bits are divided into groups of size two(8 bits). Then alternately each group selects elements from the first half of first and second vector, and then from the second vector. i.e. first group selects from first half of vec1, second group from first half of vec2, first group from second half of vec1, second group from second half of vec2.
4. **Blend:** The blend instruction can be used to select one of a specific value from one of the two registers depending on the control signal. Only those number of control bits are used as there are numbers of the specific type in the register.
If we were working with 64-bit double precision numbers then only the first 4 control bits would have been used in the above figure.

Figure 2: Unpack Lo-Hi

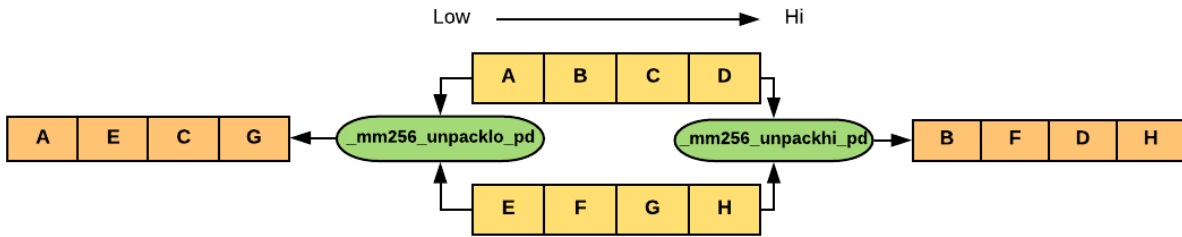


Figure 3: Shuffle

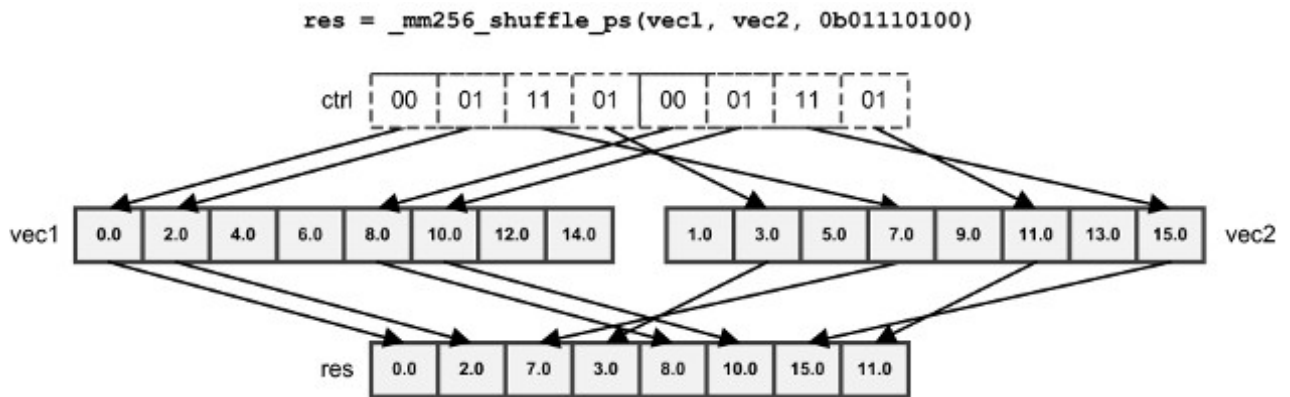



Figure 4: Blend



Its important to keep in mind that all these components have different intrinsic functions for different types and one has to be careful to use the correct type to ensure efficiency and correctness. We'll discuss this more in a later section.

3.4 SORTING NETWORKS

Sorting networks which have been well summarized in Knuth (1997) are one of the best ways to get huge speedups via SIMD. They allow for defining custom sized networks based on the AVX register size and data type available. Once the data has been loaded into registers a set of vectorized min-max operations can be run without branches and with high potential for Instruction Level Parallelism. We used this [handy utility](#) to obtain the most optimal sorting network configuration for different sized


networks. An example of the input and output to/from a 4x4 sorting network has been visualized through Figure 8.

3.5 COMPLEX SIMD OPERATIONS

In this subsection we describe more complex operations achieved using SIMD required to efficiently implement our sorting algorithm described later.

1. **Transpose:** The basic idea for transpose a matrix stored in SIMD registers is, suppose N by N matrix:
 - (a) Interleave N rows by Unit 1. (32 bits for `int`, 64 bits for `int64_t`).
 - (b) Interleave N rows by Unit 2. (64 bits for `int`, 128 bits for `int64_t`).
 - ...
 - (c) Interleave N rows by $N/2$.
 the transpose can be done in $N \log_2 N$ operations

Figure 5: 4x4 Transpose Example



2. **Horizontal Intra-Register Sort:** The idea behind intra-register sort is bitonic merging network. For a merging network, two sorted sequences of length K each are fed as inputs, and a sorted sequence of length $2K$ is produced. Figure ?? shows the idea of merging 2 sequences of 4 elements. Initially A and B are held in SIMD register and are both sorted in the ascending order. Bitonic merge requires one of the sequences sorted in descending order, so a reverse of B is done on register. Each level of the network use **Min/Max** to do comparison, after this the data in register is permuted for the appropriate comparison for the next level.

Figure 6: Bitonic Merge Network

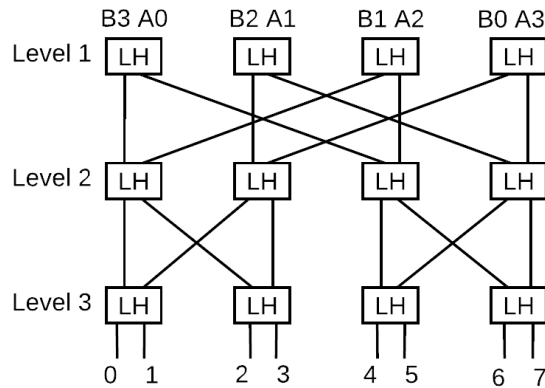


Figure 7: Intra-Register Sort for 2 Registers of 16 Values Each

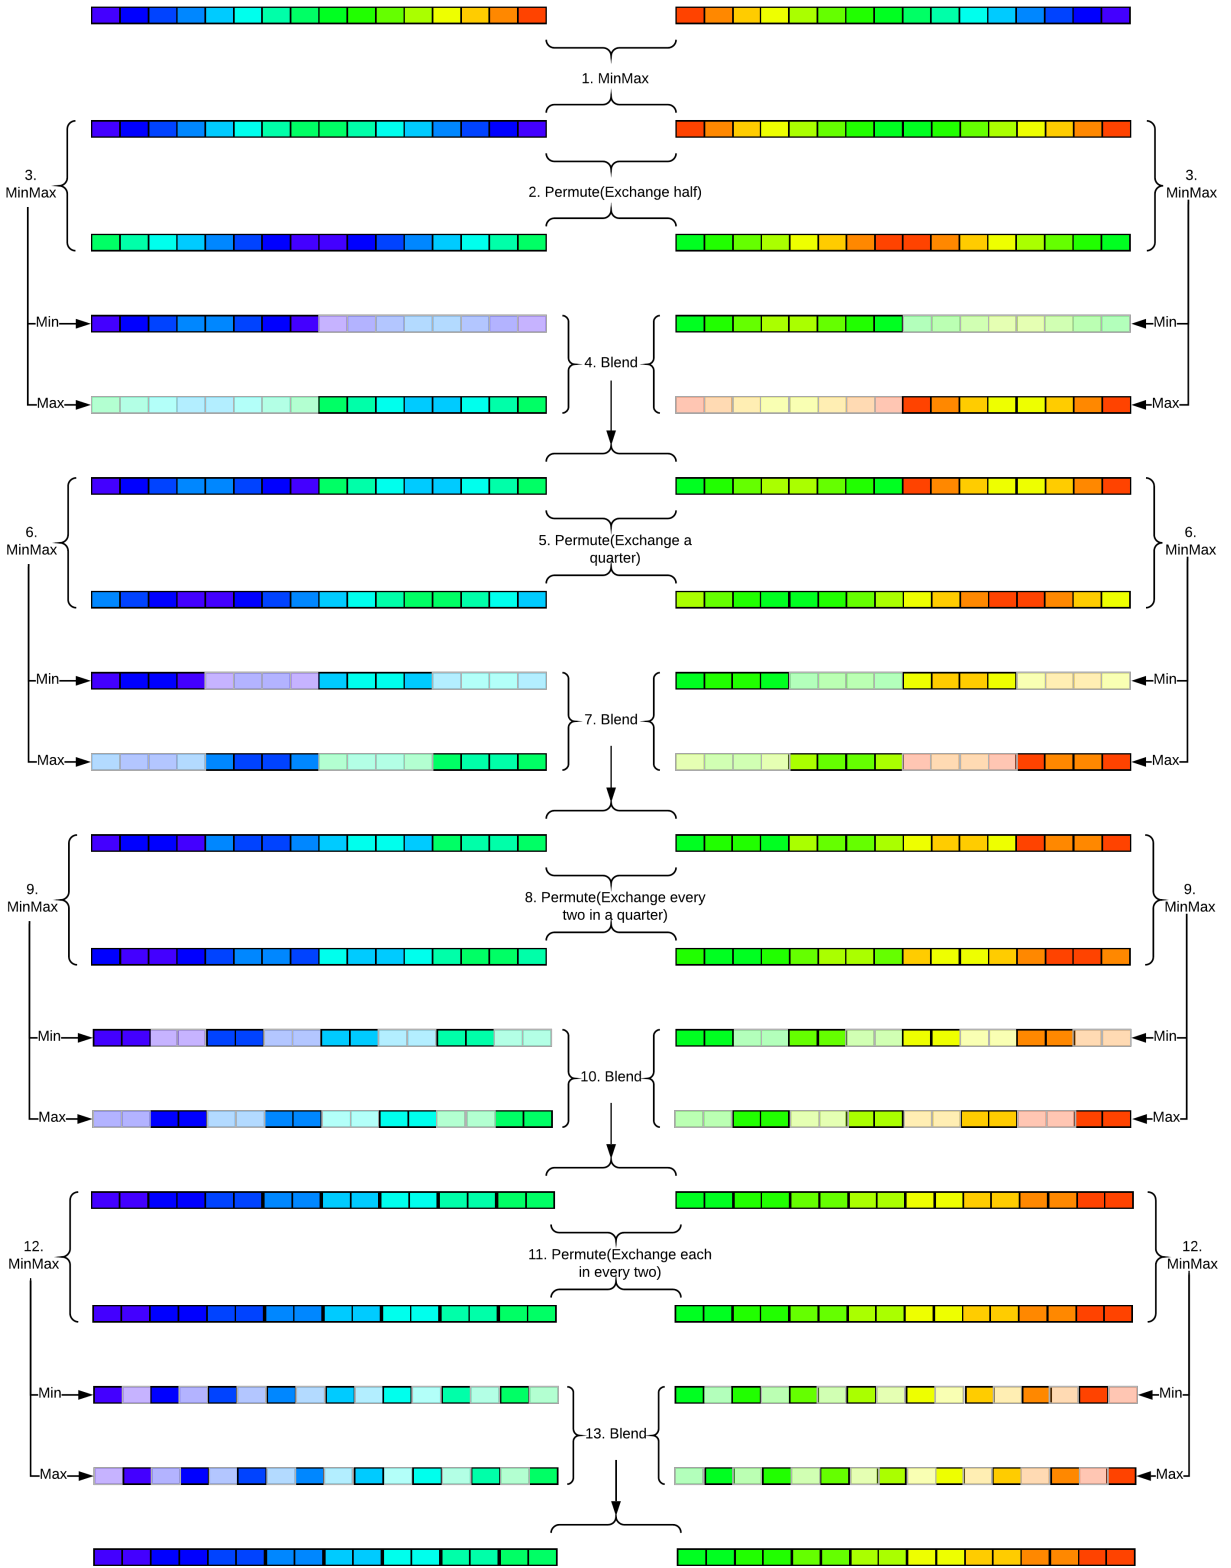


Figure ?? showed an example of merging two 16-value registers using SIMD instructions. With Blue box indicating smaller values to Red box indicating larger values, shallow color boxes are redundant duplicated values resulted from permutation. The register on the right is already flipped, the output is a sorted 32-value sequence stored in two registers respectively, with the smallest 16 values in the left register and the largest 16 values in the right register. The steps are also shown in the example to indicate the dependencies. Steps which don't have any dependencies can be executed parallelly.

4 APPROACH

In this section we will describe the main approach including the algorithms used in each step, its variants and along the way we motivated the need for a **DSL for SIMD Sorting**.

4.1 ALGORITHMS

We summarize the steps of sorting an N elements sequence here, in the subsections we will introduce our building blocks with two phases. To sort N elements, for example, consider N `int` values, each value is of 32 bits, in AVX512, one register can represent 16 such values. We pack 16 `int` values into one register, and prepare 16 such registers to form a $16 \times 16 = 256$ values, we call this a block, and then apply a 16 by 16 sorting network to this block to get 16 column-wise(vertically) sorted registers. Once we transpose these 16 registers, we can get the new 16 registers in which the 16 values are sorted. This process is done repeatedly on the N elements, 256 values at a time, with every 16 values sorted.

Next, to get a completely sorted array, a merging process is performed for those 16-value sorted sequences. This process is highly iterative, for example, the first iteration for the example above would be merging every two 16-value sorted sequences to get a 32-value sorted sequence. Next iteration would be merging every two 32-value sorted sequences to a 64-value sorted sequence, algorithm used for using 16-value register to merge 32-value sequence is detailed in the subsection. Such and so on until the whole sequence is merged. This process requires N to be the power of 2.

4.1.1 PHASE 1: BUILDING SORTED RUNS

The first step in our algorithm is to obtain sorted runs of fixed size on which the next phase will operate. In order to obtain such sorted runs we can apply the following steps:

1. Based on the data type to sort, a custom best-case sorting network has to be selected(For example for `Int32` on AVX2, we can select a network of 8×8 in the best case.) We then apply this sorting network on a block of numbers(eg. 64 in the example from above). This has been visualized in Figure 8.
2. When applied on a block of numbers the sorting network returns the numbers sorted but between same positions of different registers - i.e. vertically not horizontally sorted. Thus we need a Transpose operation - This has again been implemented as described in Section III. It's again important to keep in mind that depending on the sorting network size the correct transpose operation should be performed.


At this stage we have a huge array with sorted runs of `unit run size`(which in the example above is 4).

4.1.2 PHASE 2: MERGING SORTED RUNS

After completing phase 1 we will have an array with sorted runs. The next step would be to efficiently merge sorted runs. In the explanations that follow, `unit run size` refers to the unit length on which an operation can apply. For example the `BitonicMerge8` merges 2 registers of size 8, so the `unit run size` is 8. The algorithm we follow is similar to merge-sort and has the following building blocks:

1. **Merge-Run-Pass Block:** In this block, the `BitonicMerge` which we defined earlier is used as a kernel on runs of variable size(multiple of the `unit run size`). A high

Figure 8: Sorting Network In/Out Example



level overview of this operation is available in figure 9 and has been described in H. Inoue (2015). In the figure, we use the `BitonicMerge4` as a kernel. We start by loading two unit size runs into registers. Then at each step we apply the bitonic merge kernel, and store the smaller output to an output buffer. We then load the next sorted run in. This is continuously done until all the runs have been processed. In practice this is implemented slightly differently:

- (a) Split the input set of runs into two streams A and B.
- (b) Load R_a and R_b with the first run from A and B.
- (c) If $R_a[0]$ is less than $R_b[0]$ we select the next run to be loaded from A, else B.
- (d) Apply bitonic merge and store the smaller half out and larger half in R_b .
- (e) Load the next run in (from (c)).
- (f) Go to (c). Continue until either of the streams has been completely processed. Then just merge in any remaining portions of the stream still having unprocessed runs.

The final output has chunks of merged/sorted runs of $2 \times \text{run_size}$.

2. **End-To-End Merging:** We can now treat the `Merge-Run-Pass-Block` as a kernel itself to get finally merged runs. This is done by starting to merge unit run size of chunks of sorted runs and doubling run sizes to be merged from there onwards. Thus $\log(\text{seq_length})$ passes would be needed to complete the entire sorting procedure. This has been demonstrated in Figure 9 and has been derived from the work of C. Kim & Dubey (2009). For example if we used a sorting network of size 8×8 we will have a unit run size of 8. In the first pass we will merge 2 chunks of size 8 to merged-sorted runs of size 16. In the second pass we would merge chunks of length 32. This would go on till we reach the length of the sequence.

4.2 SORTING KEY-VALUE PAIRS

A widely used application of sorting algorithms is to sort key-value pairs. In this subsection we addressed the methodology we used to achieve this objective. When we say "key-value sort", we mean that we want to sort the values by the key. Following are the key ideas:

1. **Keep Key-Value Together:** Amongst the papers we read as part of our Literature survey, we only saw Furtak et al. (2007) which has specifically proposed an algorithm for KV pair sorting. However he recommends keeping the key-value pairs in separate registers - working on the keys and mimicing the same movements on the values. Such a design allows for little code reuse and would need complete rewriting of all the components for KV sorting. We instead propose a KV sorting design in which KV pairs are packed into

Figure 9: Merging Sorted Chunks

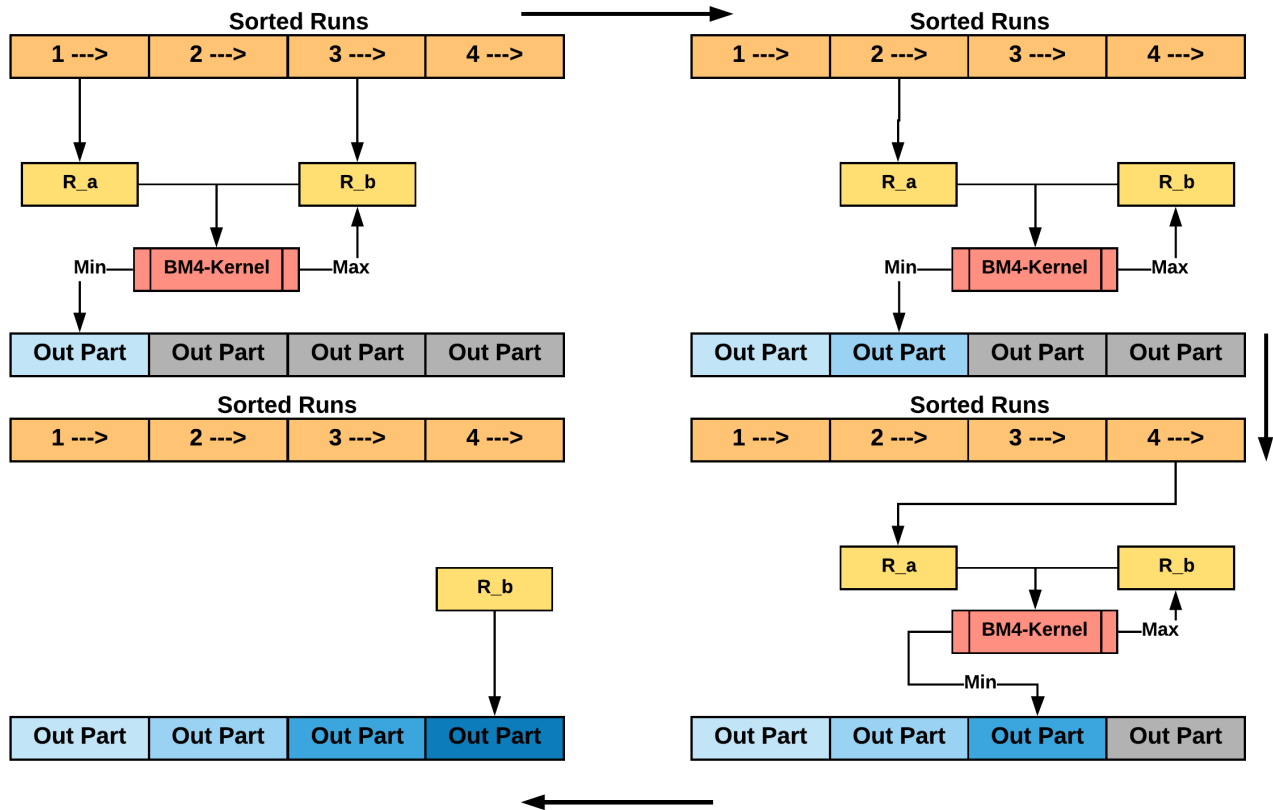
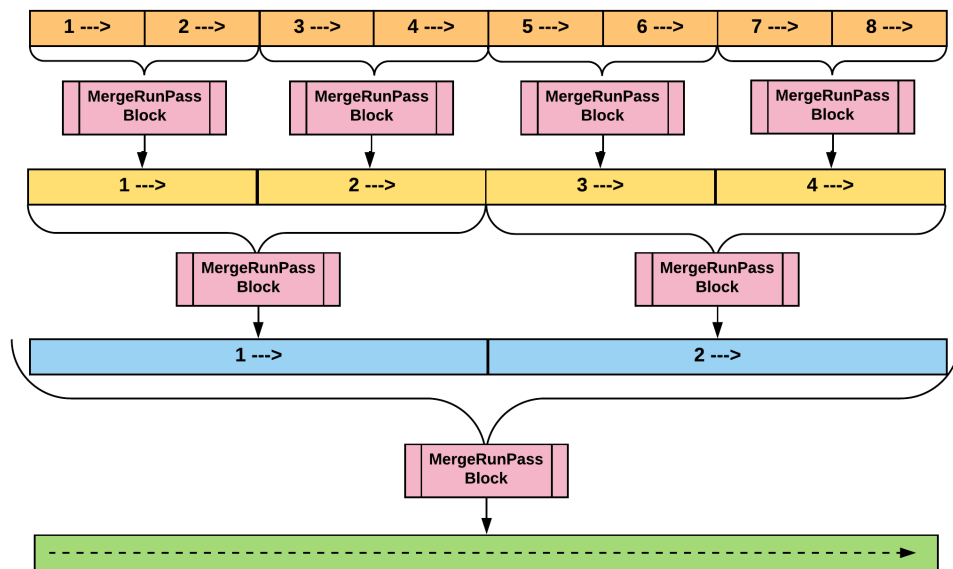


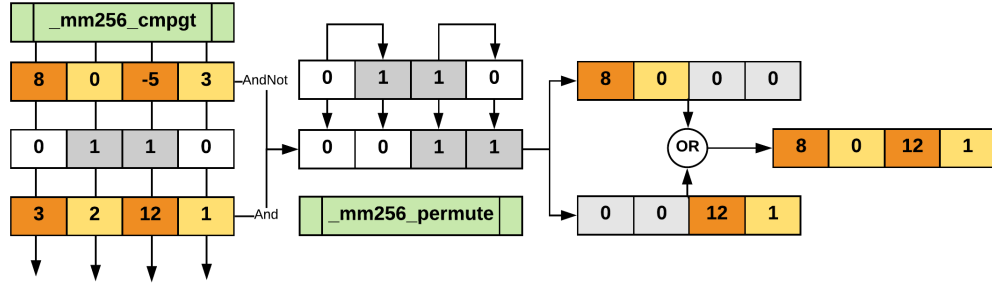
Figure 10: End-To-End Merging



the same SIMD register contiguously. Once this is done it can be noted that majority of the original code(especially the utility operations described in Section III) and design can be reused with carefully masking any comparison operations. It's important to note that the effect of doing a movement operation should be the same on registers holding different types within them since movements happen across lanes of 32, 64 or more bits without caring about values. It is at the same time important to note however that performing a movement operation with the intrinsic of correct type has much lower latency in some cases compared to intrinsics of a different type.

2. **Masked Min-Max:** The procedure to perform masked comparison operations to perform the original MinMax operations while copying the movements of the value with the key are summarized in Figure 11. Once we have a `MaskedMinMax` we can begin to construct masked versions of all other operations by simply replacing the Min-Max with the Masked Min-Max variants. Additionally this is only required in steps where there is any sort of comparison.
3. **Adjust Movement operations:** We also need to replace the original movement operations on units of same run size with operations working on double the lane size and half the data. eg. Transpose 4x4 on 64-bit values to Transpose 2x2 with 128-bit values when sorting in Phase 1.

Figure 11: Masked Min-Max



4.3 "ORDER BY" IN DBMS

The operation of sorting key-value pairs can be used as a lower-level API to perform a Database order by(albeit limited to 1 column - We weren't able to come up with an efficient generalized algorithm for multicolumn order by.) This can be done by the following steps:

1. **Convert to KV pairs:** The specific column on which we want to `Order By` goes into the `Key` field and we assign the `Record-Id` or `0-offset` in the value.
2. **KV Pair Sort:** Now apply the algorithm described in previous sections for KV sorting.
3. **Rearrange Records:** Once we have our sorted KV pairs we can use the `0-offsets` to rearrange records to their correct positions. Its important to note that this step will have a huge penalty as it will involve a large number of random accesses throughout main-memory. H. Inoue (2015) has proposed a variation of this step where we perform such rearrangements incrementally while running the algorithm. This is a possible future work for us.

4.4 PARALLELIZING SIMD SORTING

We used OpenMP to try and parallelize our code. C. Kim & Dubey (2009) have done an excellent job in this aspect by allowing for parallelized merge pass, multiway merges and carefully distributed data placement, we attempted to parallelize Phase 1 of the algorithm and the merge pass for lack of time. Following is our methodology and inferences:

1. **Parallelizing creation of Sorted Runs:** The sorted runs are simple to parallelize since they involve independent application of the Bitonic Sorting network on chunks of fixed size. Unfortunately this step is highly bandwidth bound. Thus instead of seeing a speedup with more threads we saw a drop in speed because of the additional overhead of fork-join and task assignment with no gains in terms of latency due to being bandwidth bound.
2. **Parallelizing Merge Pass:** As can be seen in Figure 10, the merge pass involves application of the `MergeRunPass` kernel on chunks of equal size. This can be parallelized as long as we ensure that all relevant variables and the buffer offset is correctly setup. When we applied this, we did find a speedup, but this was sub-linear. This again is because there is a lot of load-store operations from arrays into registers and vice versa during operations.

4.5 DSL SETUP AND GENERALIZING SIMD SORTING

As we’ve seen in the above subsections, there is a strong motivation to build a DSL for SIMD sorting:

1. Between register types the nature of supported functions widely varies. Thus what is most optimal strongly depends on the sorting problem and hardware in question.
2. Different CPUs support different types of registers: SSE4, AVX, AVX2, AVX512. We always want to use the most optimal register type available on the hardware in question.
3. Different functions within a register for different data types: Thus most available implementations only support integer data types. However this is not true of real world workloads which have varying types.

Thus, we propose **UltraSort** as a DSL for SIMD sorting which:

- Work across different hardware architectures(supporting and optimizing differently for both AVX2 and AVX512.)
- Support most optimized variants for every numeric data type. Based on the data type used the most optimal sorting configuration possible is picked.
- Support key-value pair sorting for any numeric data type.
- Supports record-level sorting based on a single "key".
- Parallel execution(currently in a multicore non-distributed manner).

As can be seen in Figure 12 our DSL framework chooses the optimal SIMD sorting variant based on different characteristics of the workload and hardware. At a higher level is the selection of the optimal kernels and then at a lower level the correct set of functions(by data type) for the implementation of the kernel. For example in the case of Bitonic sorting for int32 on AVX512 we can have a 16x16 bitonic sorting network possible but only 8x8 on AVX2. Further the implementation varies since there are different function signatures and behaviors on both architectures.

Wherever possible we’ve tested for and templated functions and kernels, performing register type-casting only where there were no penalties(`_mm_cast` has zero latency but certain movement operations on a different data type than expected incur additional overheads). This also allowed us to have extensive code reuse.

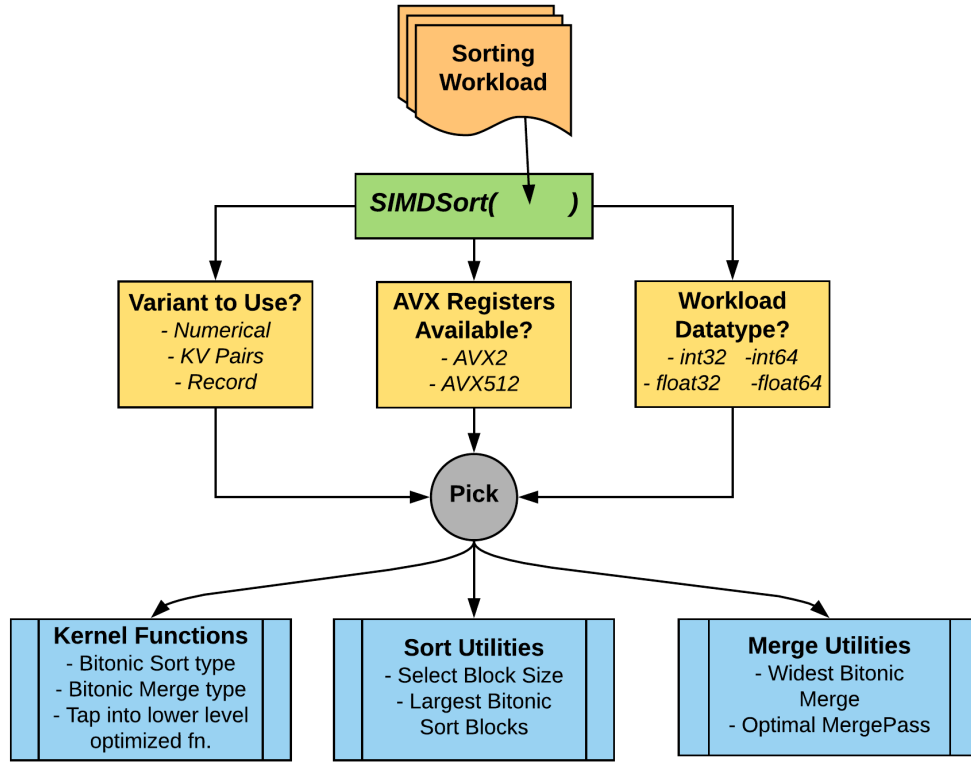
4.6 TESTING AND CORRECTNESS

We’ve tested every component of the entire framework from lower level functions to higher level kernels to the entire sort correctness within the **GTest testing framework**. This allowed us to confidently play around with changes without worrying about breaking anything. This was extremely crucial given that our complete project stands at **4000** lines of code.(Just FYI our tests also cover **4000** lines of code).

5 BENCHMARKING & RESULTS

We obtained the results from all settings listed in our proposal, with `OpenMP` number of threads set to 1. Due to large experiments results, in this section, representative results would be presented. The

Figure 12: DSL Framework for SIMD Sorting



full timing and speedup results can be found at Section 8. Table 1, 2, 3 showed the time in seconds, ipc and speedup of those different algorithms took to sort 1M elements of different settings.

Table 1: 1M Elements Sort Time(in seconds)

Settings/Algorithms	std::stable_sort	std::sort	ips4o::sort	pdqsort	avx256::sort	avx512::sort
32BitInteger	0.063970	0.080234	0.035442	0.036089	0.019070	0.010364
32BitFloat	0.080819	0.090239	0.037985	0.041248	0.024958	0.020181
64BitInteger	0.065461	0.082378	0.037092	0.036340	0.028105	0.028105
64BitFloat	0.083170	0.091735	0.040148	0.041677	0.045282	0.036061
32BitKeyValueInt	0.099336	0.082743	0.036834	0.083273	0.068085	0.033655
64BitKeyValueInt	0.115010	0.085000	0.044091	0.083327	0.109607	0.106933
32BitKeyValueFloat	0.113443	0.097009	0.042667	0.097496	0.075137	0.061885
64BitKeyValueFloat	0.128100	0.099286	0.049526	0.098907	0.116892	0.107151

Table 2: 1M Elements Sort IPC(Instructions Per Cycle)

Settings/Algorithms	std::stable_sort	std::sort	ips4o::sort	pdqsort	avx256::sort	avx512::sort
32BitInteger	1.506790	0.681694	2.076354	2.336017	2.152515	1.630018
32BitFloat	1.197460	0.613799	1.823862	2.090052	1.478475	1.060270
64BitInteger	1.457014	0.670313	1.999275	2.285352	1.416596	1.157502
64BitFloat	1.169164	0.623366	1.794311	2.050244	1.201110	1.065701
32BitKeyValueInt	0.961099	0.789943	2.395018	0.612214	1.292443	1.060222
64BitKeyValueInt	0.900830	0.767878	2.345641	0.602405	1.295691	1.335133
32BitKeyValueFloat	0.861952	0.697731	2.085952	0.581825	1.066630	1.504925
64BitKeyValueFloat	0.797310	0.683473	2.090357	0.581601	1.194682	1.348604

Table 3: 1M Elements Sort Speedup(Relative to `std::sort`)

Settings/Algorithms	<code>std::stable_sort</code>	<code>std::sort</code>	<code>ips4o::sort</code>	<code>pdqsort</code>	<code>avx256::sort</code>	<code>avx512::sort</code>
32BitInteger	1.254256	1.000000	2.263849	2.223261	4.200139	7.741558
32BitFloat	1.116557	1.000000	2.375618	2.187715	3.594582	4.471524
64BitInteger	1.258417	1.000000	2.220918	2.266862	1.484623	2.931080
64BitFloat	1.102980	1.000000	2.284896	2.201074	2.043218	2.543850
32BitKeyValueInt	0.832961	1.000000	2.246350	0.993639	1.230044	2.458543
64BitKeyValueInt	0.739063	1.000000	1.927807	1.020069	0.773838	0.794888
32BitKeyValueFloat	0.855136	1.000000	2.273620	0.995006	1.302896	1.567567
64BitKeyValueFloat	0.775070	1.000000	2.004722	1.003840	0.856764	0.926604

Table 4, 5, 6 showed the time in seconds, ipc and speedup of those different algorithms took to sort 1024 elements of different settings.

Table 4: 1024 Elements Sort Time(in seconds)

Settings/Algorithms	<code>std::stable_sort</code>	<code>std::sort</code>	<code>ips4o::sort</code>	<code>pdqsort</code>	<code>avx256::sort</code>	<code>avx512::sort</code>
32BitInteger	0.000050	0.000056	0.000098	0.000038	0.000026	0.000010
32BitFloat	0.000061	0.000063	0.000042	0.000041	0.000029	0.000012
64BitInteger	0.000069	0.000081	0.000057	0.000051	0.000047	0.000024
64BitFloat	0.000114	0.000063	0.000061	0.000039	0.000045	0.000032
32BitKeyValueInt	0.000095	0.000083	0.000058	0.000077	0.000056	0.000024
64BitKeyValueInt	0.000069	0.000060	0.000138	0.000056	0.000056	0.000052
32BitKeyValueFloat	0.000076	0.000068	0.000085	0.000061	0.000044	0.000035
64BitKeyValueFloat	0.000075	0.000066	0.000087	0.000062	0.000062	0.000052

Table 5: 1024 Elements Sort IPC(Instructions Per Cycle)

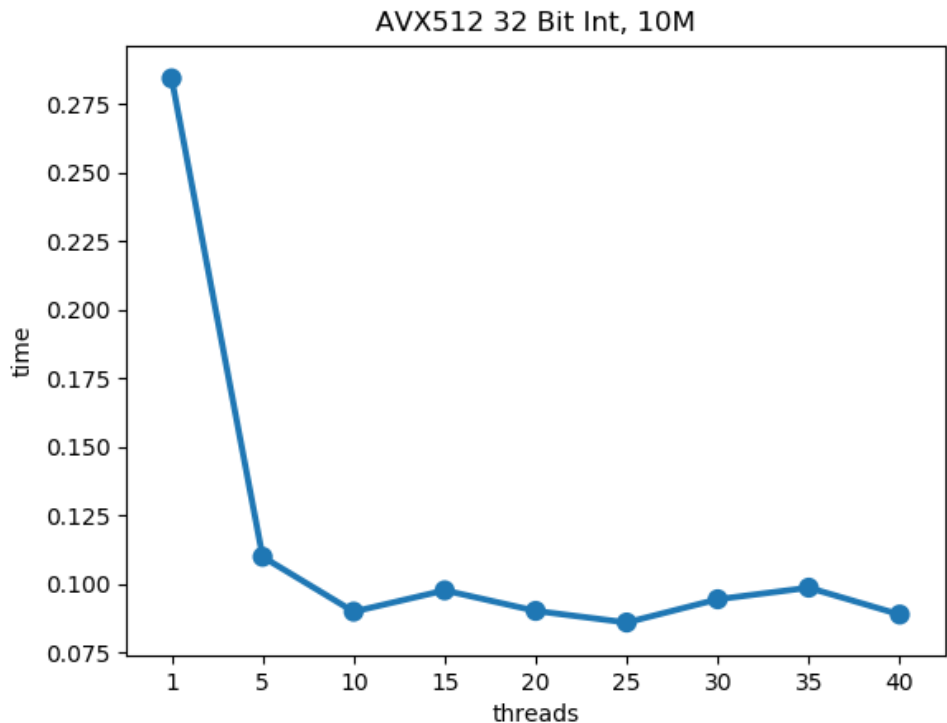
Settings/Algorithms	<code>std::stable_sort</code>	<code>std::sort</code>	<code>ips4o::sort</code>	<code>pdqsort</code>	<code>avx256::sort</code>	<code>avx512::sort</code>
32BitInteger	1.506790	0.681694	2.076354	2.336017	2.152515	1.630018
32BitFloat	1.197460	0.613799	1.823862	2.090052	1.478475	1.060270
64BitInteger	1.457014	0.670313	1.999275	2.285352	1.416596	1.157502
64BitFloat	1.169164	0.623366	1.794311	2.050244	1.201110	1.065701
32BitKeyValueInt	0.961099	0.789943	2.395018	0.612214	1.292443	1.060222
64BitKeyValueInt	0.900830	0.767878	2.345641	0.602405	1.295691	1.335133
32BitKeyValueFloat	0.861952	0.697731	2.085952	0.581825	1.066630	1.504925
64BitKeyValueFloat	0.797310	0.683473	2.090357	0.581601	1.194682	1.348604

Table 6: 1024 Elements Sort Speedup(Relative to `std::sort`)

Settings/Algorithms	<code>std::stable_sort</code>	<code>std::sort</code>	<code>ips4o::sort</code>	<code>pdqsort</code>	<code>avx256::sort</code>	<code>avx512::sort</code>
32BitInteger	1.137291	1.000000	0.575207	1.500533	3.004192	5.689899
32BitFloat	1.047894	1.000000	1.500710	1.561654	2.191511	5.296327
64BitInteger	1.178717	1.000000	1.426102	1.579914	1.693648	3.432088
64BitFloat	0.553212	1.000000	1.037420	1.633755	1.464924	1.967933
32BitKeyValueInt	0.872149	1.000000	1.428250	1.067589	1.470357	3.387044
64BitKeyValueInt	0.866976	1.000000	0.431514	1.067979	1.055197	1.154962
32BitKeyValueFloat	0.900092	1.000000	0.802920	1.113469	1.478458	1.955836
64BitKeyValueFloat	0.876327	1.000000	0.755606	1.059182	1.061596	1.281335

We also tested using more than 1 threads in AVX512 to experiment. Figure 13 shows the trend of time taken for `avx512::sort` under different number of threads setting, N is 10 million `int` data. The parallelism we took is for the phase 2, using **OpenMP** to parallelize the merge pass process.

Figure 13: Number of threads and Time in seconds of `avx512::sort`




6 ANALYSIS

6.1 SEQUENCE PERFORMANCE

To better view the statistics we collected in Section 5, we plotted them out. Since AVX256 version is only for reference so only AVX512 is plotted. We can clearly see from Figure ?? that for sorting

Figure 14: Algorithms Comparison on Different Settings, 1024 Elements on the Left, 1M on the Right



1024 elements, `avx512::sort` beats all other sorting algorithms across all different settings. Due to native instructional support for 32 bits `int` and `float` value, it achieved the best timing. Compared with Figure ??,

Dee: Analysis of Seq performance between different algorithms - Where SIMD works best - Where not and why on all applications(Single number, pairs, Order by)

Dee: Analysis of Seq performance in terms of IPC)

Dee: AVX2 vs. AVX512 performance.

Dee: KV Pair sorting vs. Record sorting - Overheads - Comparison against other sorting algos.

Dee: Final Summary - Which types of settings does SIMD sorting make sense?

7 CONCLUSION

Dee: Put conclusion

REFERENCES

J. Teubner M. Tamer Ozsu C. Balkesen, G. Alonso. Multi-core, main-memory joins: Sort vs. hash revisited. *PVLDB*(1), 7, 2013.

J. Chhugani T. Kaldewey A. D. Nguyen A. D. Blas V. W. Lee N. Satish C. Kim, E. Sedlar and P. Dubey. Sort vs. hash revisited: Fast join implementation on modern multi-core cpus. *PVLDB*2(2), pp. 1378–1389, 2009.

Timothy Furtak, José Nelson Amaral, and Robert Niewiadomski. Using simd registers and instructions to enable instruction-level parallelism in sorting algorithms. In *Proceedings of the nineteenth annual ACM symposium on Parallel algorithms and architectures*, pp. 348–357. ACM, 2007.

K. Taura H. Inoue. Simd- and cache-friendly algorithm for sorting an array of structures. *PVLDB*, 8, 2015.

Donald Ervin Knuth. *The art of computer programming*, volume 3. Pearson Education, 1997.

8 APPENDIX

8.1 EXPERIMENT MACHINE STATISTICS

8.2 RAW RESULTS